

Exercise 9: AMQ 7.12 Message Persistence & Replication

AMQ7.12 has various message persistence options:

- NIO (Java NIO, journal-based persistence)
- ASYNCIO (Linux Asynchronous IO, journal-based persistence)
- JDBC (persist to the relational database of your choice)
- Memory (in-memory stateless message persistence)

The default, out-of-the-box setting is NIO file journal-based persistence. You can configure message persistence by updating the `/jboss-amq-7.0.0.redhat-1/brokers/myfirstbroker/etc/broker.xml` file. Here is the section in `broker.xml`:

```
<configuration>
  <core>
    ...
    <persistence-enabled>true</persistence-enabled>
    <!-- this could be ASYNCIO or NIO-->
    <journal-type>NIO</journal-type>
    <paging-directory>./data/paging</paging-directory>
    <bindings-directory>./data/bindings</bindings-directory>
    <journal-directory>./data/journal</journal-directory>
    <large-messages-directory>./data/large-messages</large-messages-directory>
    <journal-datasync>true</journal-datasync>
    <journal-min-files>2</journal-min-files>
    <journal-pool-files>-1</journal-pool-files>
    ...
  </core>
</configuration>
```

Replication

The following lab demonstrates setting up a 2-node master / slave cluster with shared-nothing replication.

Lab - Shared-nothing replication between Master / Slave cluster

Prerequisites

Download the legacy ActiveMQ client JAR which supports OpenWire failover [here](#)

1. Create a master and slave broker pair by running the following commands:

```
$ ./bin/artemis create brokers/master
$ ./bin/artemis create brokers/slave --port-offset 1
```

Be sure to give each broker the `admin/admin` username / password combo and allow anonymous access.

2. Replace the `etc/broker.xml` file from the examples directory for both master and slave:

```
cp /jboss-amq-7.0.0.redhat-1/examples/features/clustered/clustered-static-  
discovery/src/main/resources/activemq/server0/broker.xml  
/brokers/master/etc/broker.xml  
cp /jboss-amq-7.0.0.redhat-1/examples/features/clustered/clustered-static-  
discovery/src/main/resources/activemq/server1/broker.xml  
/brokers/slave/etc/broker.xml
```

3. Replace the security-settings section of both master / slave `broker.xml` files with the following text (this will allow the producer / consumer to dynamically create queues):

```
<security-settings>  
  <security-setting match="#">  
    <permission type="createNonDurableQueue" roles="amq"/>  
    <permission type="deleteNonDurableQueue" roles="amq"/>  
    <permission type="createDurableQueue" roles="amq"/>  
    <permission type="deleteDurableQueue" roles="amq"/>  
    <permission type="createAddress" roles="amq"/>  
    <permission type="deleteAddress" roles="amq"/>  
    <permission type="consume" roles="amq"/>  
    <permission type="browse" roles="amq"/>  
    <permission type="send" roles="amq"/>  
    <!-- we need this otherwise ./artemis data imp wouldn't work -->  
    <permission type="manage" roles="amq"/>  
  </security-setting>  
</security-settings>
```

4. Update the cluster-connection section by replacing the message-load-balancing line with the following:

```
<message-load-balancing>ON_DEMAND</message-load-balancing>
```

This will prevent message starvation and enable message redistribution between nodes.

5. Startup both the master and slave brokers in separate consoles

```
./brokers/master/bin/artemis run  
./brokers/slave/bin/artemis run
```

6. Using the legacy `activemq` client, run the following commands in two separate console windows:

```
java -jar activemq-all-5.11.0.redhat-630187.jar consumer --brokerUrl  
'failover:(tcp://localhost:61616,tcp://localhost:61617)' --user admin --  
password admin --destination queue://TEST  
  
java -jar activemq-all-5.11.0.redhat-630187.jar producer --sleep 100 --  
messageCount 1000 --user admin --password admin --brokerUrl
```

```
'failover:(tcp://localhost:61616,tcp://localhost:61617)' -destination  
queue://TEST
```

7. Kill the master broker, and observe failover of both consumer / producer processes to the slave broker
8. Startup the original master broker again. Kill the slave broker, and notice failover back to the original master.